

# Tutorial: Manage Apps and Cloud Resources in Unified Approach with Kubernetes

*Jianbo Sun, Andy Shi @Alibaba & Jared Watts @Upbound*



# Speakers



KubeCon



CloudNativeCon

North America 2020

*Virtual*




**Jianbo Sun**

Senior Engineer @Alibaba, Cloud Native App Platform team, working on Alibaba's web scale app delivery system and OAM/KubeVela



**Andy Shi**

Staff Engineer,  
Alibaba Cloud



**Jared Watts**

Co-creator of @crossplane\_io  
Founding Engineer @upbound\_io

# Checkpoints



KubeCon



CloudNativeCon

North America 2020

*Virtual*

1. Prerequisites
2. Introducing KubeVela
3. Ship the first cloud native application with KubeVela
4. Introducing Crossplane
5. Prepare cloud resources with Crossplane
6. Configure the application to consume cloud resources
7. Ship the application to another cloud w/o modification
8. Wrap up

# Prerequisites



KubeCon



CloudNativeCon

North America 2020

*Virtual*

Instruction and files at

<https://github.com/oam-dev/kubevela/tree/master/documentation/kubecondemo>

1. A Kubernetes cluster >1.16
2. A cloud provider api key and secret
3. Download vela following  
<https://github.com/oam-dev/kubevela/blob/master/README.md>
4. Install crossplane (later)
5. Install Kubevela
6. `sudo cp bin/vela /usr/local/bin/vela`
7. `vela install`

# Who Are We?



KubeCon



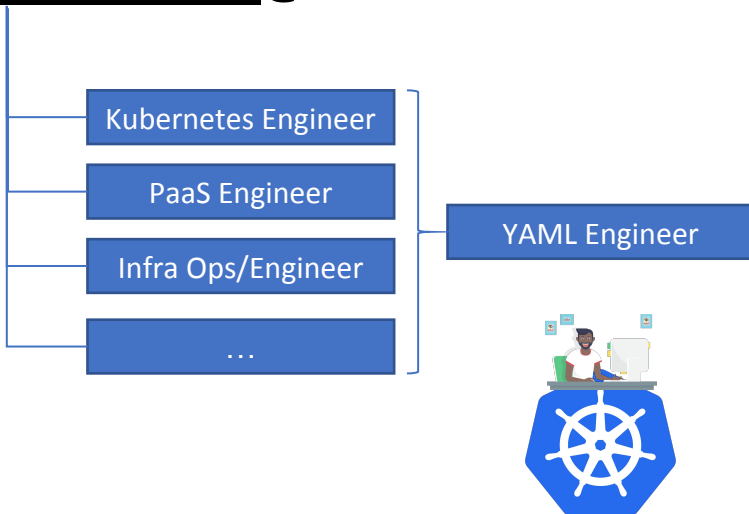
CloudNativeCon

North America 2020

Virtual

We are:

- **Platform Builders** @Alibaba



## How Does Alibaba Ensure the Performance of System Components in a 10,000-node Kubernetes Cluster?

Alibaba Developer    October 24, 2019    👁 23,653    💬 0

This article looks the problems and challenges that Alibaba Cloud overcame for Kubernetes to function at an ultra-large scale and the specific solutions proposed.

*By Zeng Fansong, senior technical expert for the Alibaba Cloud Container Platform, and Chen Jun, systems technology expert at Ant Financial.*

This article will take a look at some of the problems and challenges that Alibaba and its ecosystem partner [Ant Financial](#) had to overcome for [Kubernetes](#) to function properly at mass scale, and will cover the solutions proposed to the various problems the Alibaba engineers encountered. Some of these solutions include improvements to the underlying architecture of the Kubernetes deployment, such as enhancements to the performance and stability of etcd, the kube-apiserver, and kube-controller. These were all crucial for Alibaba to ensure the support needed for the 2019 [Tmall 618 Shopping Festival](#) to take full advantage of the 10,000-node Kubernetes cluster deployment. They are also important lessons for any enterprise interested in following Alibaba's footsteps.

# What We Build?



KubeCon

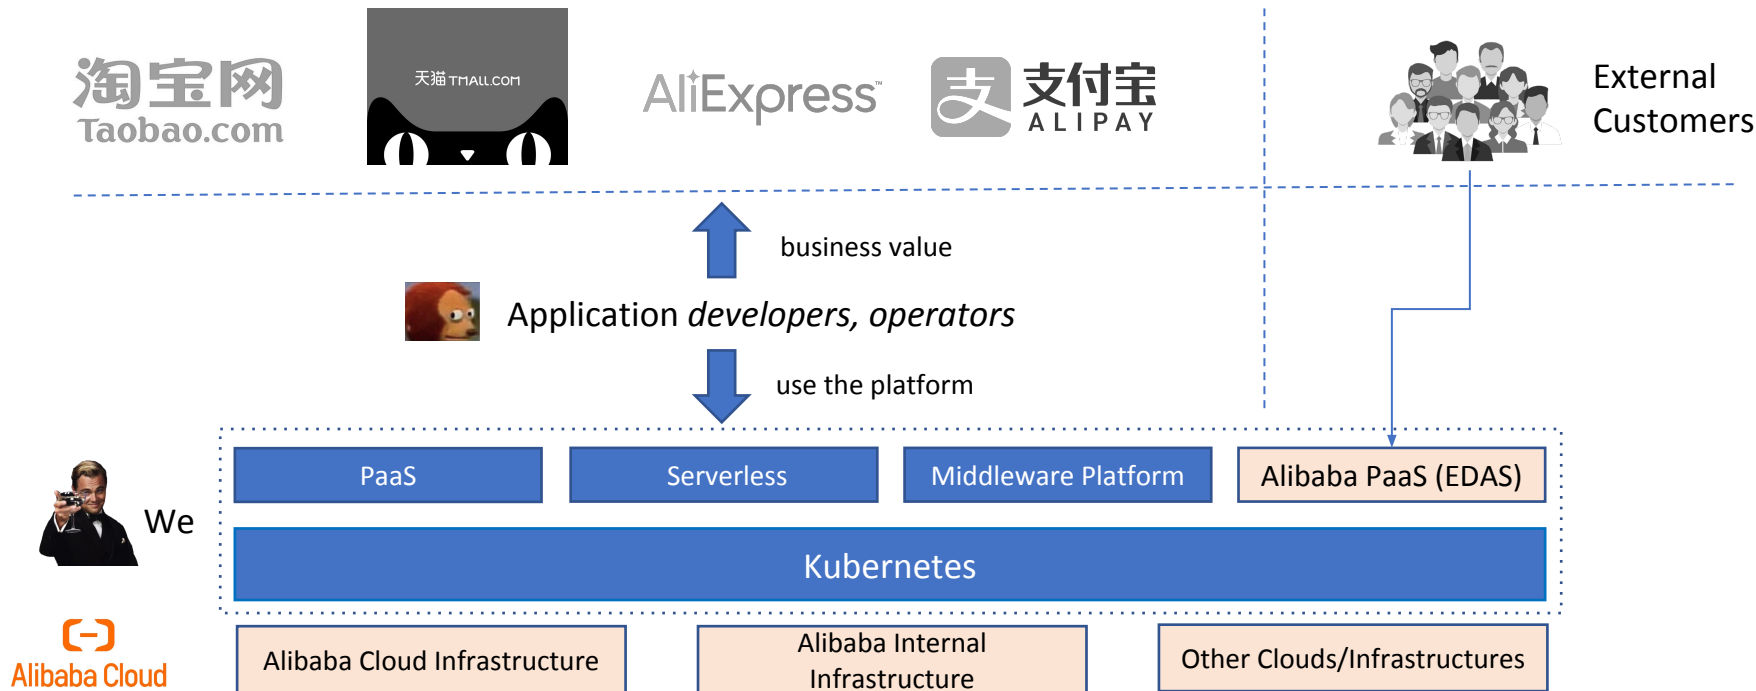


CloudNativeCon

North America 2020

Virtual

- Well ... lots of platforms on top of k8s, in hybrid environments



# Why Build App Platforms?



KubeCon



CloudNativeCon

North America 2020

Virtual


Bring **application** context back to k8s!



Application *developers, operators*



## App-Centric API



*what the platforms provide*

## App-Centric Abstractions

### scaling

- auto scale +100 instances when latency > 10%

### rollout

- promote the canary instance with step of 10%

## App-Centric User Interfaces



GUI



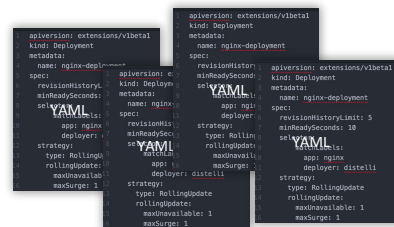
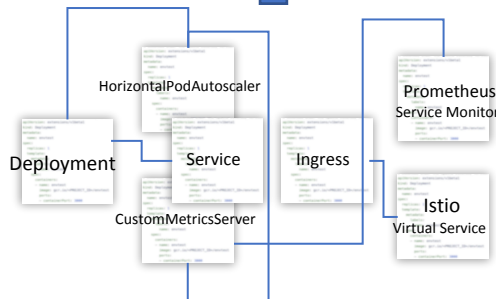

CLI



IaC



*what k8s provides*



# Why Build Unified App Platforms?



KubeCon



CloudNativeCon

North America 2020

Virtual

## Building app platforms is hard ...

- **Fragmentation:** ~11 PaaS/Serverless in Alibaba
- **Silos:** no interoperability, reusability, or portability
- **Close:** many in-house wheels due to in-house app crd

Users



I run stateful workloads!

I run stateless apps!

I run stateless serverless containers!



Platform Builders

App CRD

Canary

Manual  
Scaling

Cert

AutoScaler

Deployment

Job

Cert  
Manager

Flagger

HPA

Ingress

Virtual  
Service

Knative  
Service

Route

AutoScaling

Let's  
Encrypt

Kubernetes



# Introducing KubeVela



KubeCon

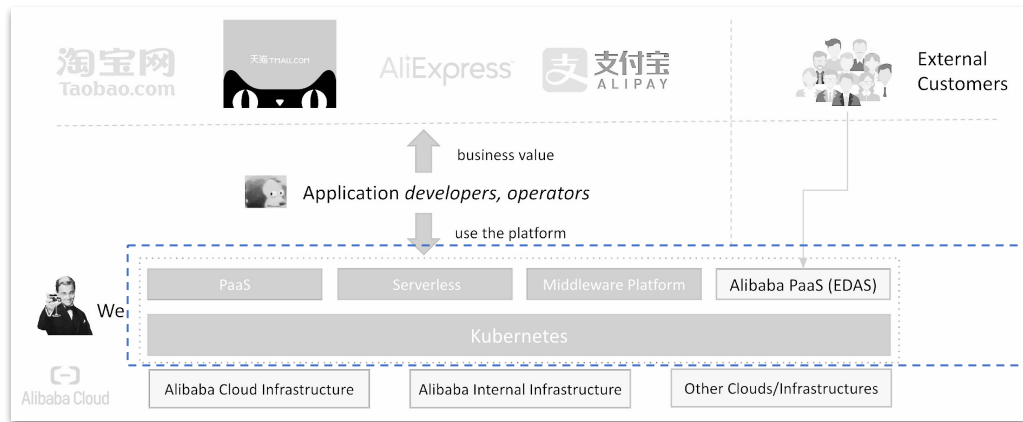


CloudNativeCon

North America 2020

Virtual

- Come from real-world practices:
  - KubeVela's design originated from Alibaba's **unified application platform engine**
- Community owned project:
  - Initialized by **bootstrapping contributors** from 8+ different organizations since day 0
  - Incubated by OAM community and intend to be donated to neutral foundation at early stage



<https://github.com/oam-dev/kubevela>

# The Goal of KubeVela



KubeCon

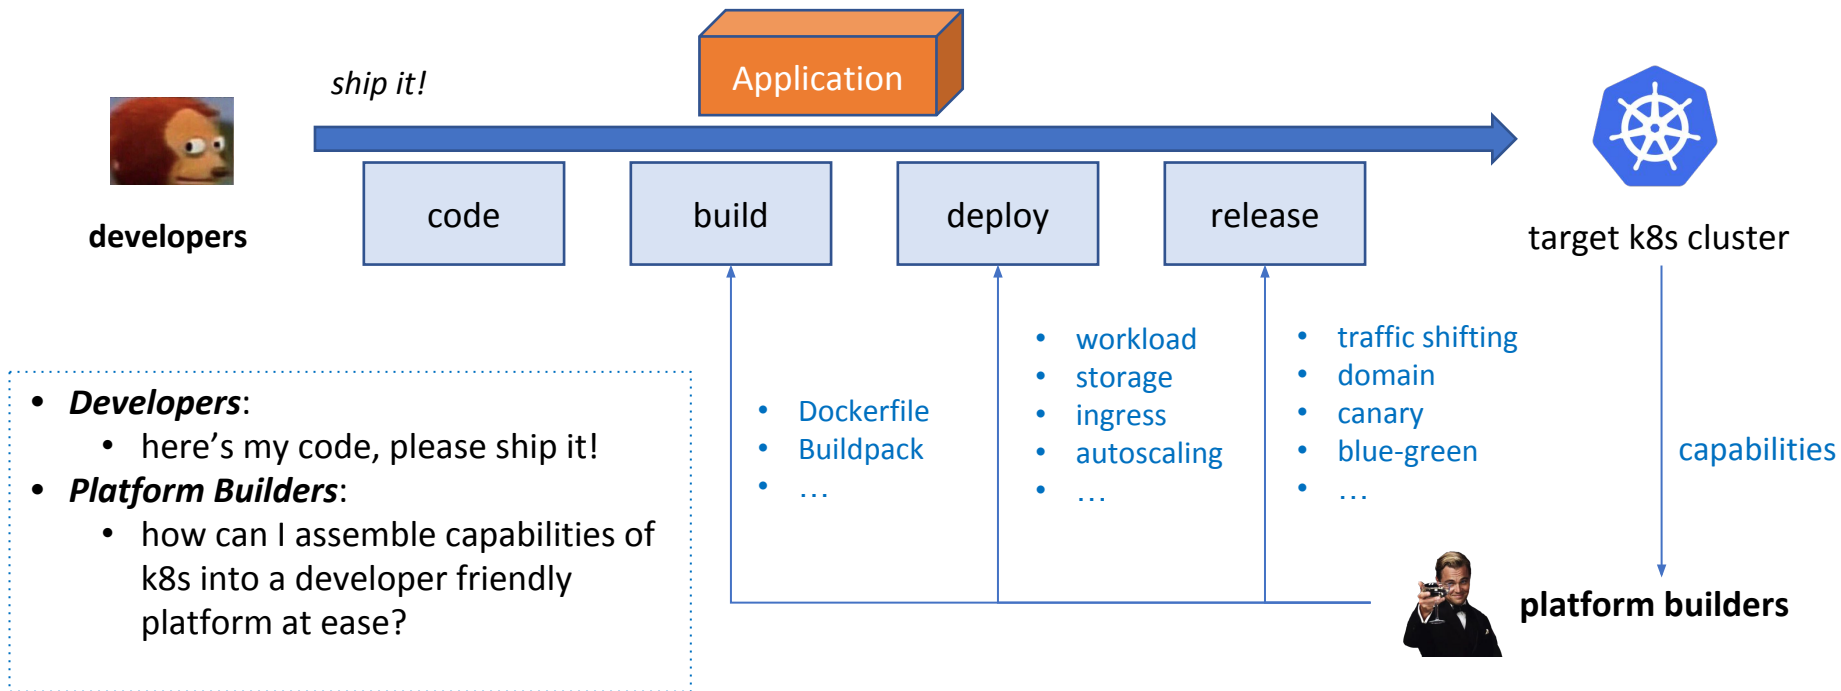


CloudNativeCon

North America 2020

Virtual

KubeVela aims at **both** *developers* and *platform builders*



- **Application-Centric**

- We believe “application” should be the main (maybe the only?) API our platform exposes to users.
- e.g. KubeVela adopts *Open Application Model (OAM)* as the app-centric API and introduces *Appfile* as the last mile developer tool

- **Capability Oriented Architecture (COA)**

- **Every** feature in KubeVela is a independent plugin (either a k8s built-in resource or your own CRD controller).
- e.g. Alibaba use KubeVela adopts *Flagger* as *rollout* trait, *KEDA* as *autoscaling* trait

- **Highly extensible, even for its user interface**

- When a new capability is installed, it should immediately consumable by end users without re-compiling or re-installing KubeVela.
  - e.g. KubeVela’s *Appfile*

# Exercise 1



KubeCon



CloudNativeCon

North America 2020

*Virtual*

## Ship the first cloud native application with KubeVela

1. `vela system update`
2. `vela workloads`
3. `vela traits`
4. `vela comp deploy mycomp -t webservice --image  
crccheck/hello-world --port 8000 --app myapp`
5. `vela app status`

# Appfile



KubeCon



CloudNativeCon

North America 2020

Virtual

```
services:
  express-server:
    build:
      image: oamdev/testapp:v1
    docker:
      file: Dockerfile
      context: .

    cmd: ["node", "server.js"]

  route:
    domain: example.com
    http: # match the longest prefix
      "/": 8080

env:
- F00=bar
- F002=sec:my-secret # map the key sam
- F003=sec:my-secret:key # map specifi
- sec:my-secret # map all KV pairs fro

files: # Mount secret as a file
- /mnt/path=sec:my-secret

scale:
  replica: 2
  auto: # automatic scale up and down ba
    range: "1-10"
  cpu: 80 # if cpu utilization is abov
  qps: 1000 # if qps is higher than 1k

canary: # Auto-create canary deployment.
  replica: 1 # canary deployment size
headers:
- "foo:bar.*"
```

## Capability Definition 1

```
apiVersion: core.oam.dev/v1alpha2
kind: WorkloadDefinition
metadata:
  name: webservice
spec:
  definitionRef:
    name: deployments.apps
  extension:
    template: |
      parameter: #webservice
      #webservice: {
        // +vela:cli:enabled=true
        // +vela:cli:usage=specify commands to run in container
        // +vela:cli:shortsc
        cmd: [...string]

        env: [...string]

        files: [...string]
      }

    output: {
      apiVersion: "apps/v1"
      kind: "Deployment"
      metadata:
        name: context.name
      spec: {
        selector: {
          matchLabels:
            app: context.name
        }
        template: {
          metadata:
            labels:
              app: context.name
          spec: {
            containers: [{
              name: context.name
              image: context.image
              command: parameter.cmd
            }]
          }
        }
      }
    }
}
```

```
apiVersion: core.oam.dev/v1alpha2
kind: TraitDefinition
metadata:
  name: route
spec:
  definitionRef:
    name: routes.standard.oam.dev
  extension:
    template: |
      parameter: #route
      #route: {
        domain: string
        http: {string}: int
      }

    // trait template can have multiple outputs and they are all traits
    outputs: service: {
      apiVersion: "v1"
      kind: "Service"
      metadata:
        name: context.name
      spec: {
        selector:
          app: context.name
        ports: [
          for k, v in parameter.http {
            port: v
            targetPort: v
          }
        ]
      }
    }
}
```

## Capability Definition 2

## • Simple

- Think about docker-compose but for Kubernetes.
- Designed to ship (build -> release) cloud native app by one click.

## • Extensible

- Every section in Appfile references a independent capability definition

## • CUE based

- The schema of each section is enforced by CUE template defines in capability definition.

# How Does Appfile Work?



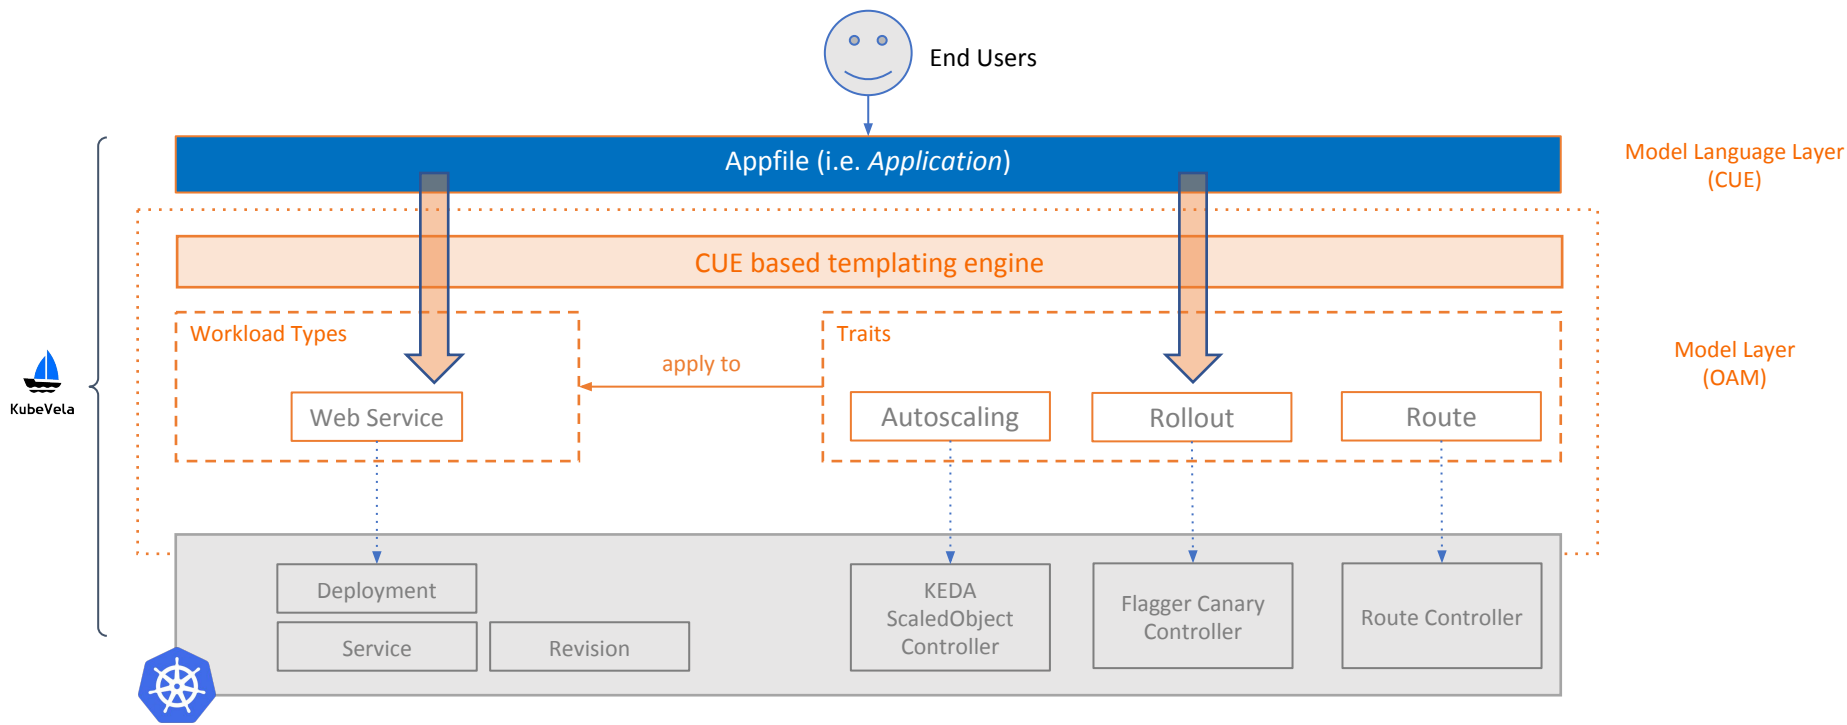
KubeCon



CloudNativeCon

North America 2020

Virtual



# Extend KubeVela!



KubeCon



CloudNativeCon

North America 2020

Virtual

Any capability in the open source community, could be shipped as a **feature** of KubeVela by a simple `$ kubectl apply -f definition.yaml`, and become usable immediately to users through *Appfile* and other UI tools

- This is also how all built-in features in KubeVela are shipped btw.

Let's say:

- add a new workload type of
  - AWS RDS?
  - stateful workload?
  - Redis Operator?
- add a new capability such as:
  - app metrics?
  - blue-green rollout?
  - custom domain?
  - ... enable Dapr?



`$ kubectl apply -f`

```
apiVersion: core.oam.dev/v1alpha2
kind: TraitDefinition
metadata:
  name: metric
  namespace: default
  annotations:
    definition.oam.dev/apiVersion: standard.oam.dev/v1alpha1
    definition.oam.dev/kind: MetricsTrait
    definition.oam.dev/description: "Add metric monitoring for workload"
spec:
  appliesToWorkloads:
    - webservice
    - backend
    - task
    - containerizedworkloads.core.oam.dev
    - clonesetworkloads.apps.kruise.io
    - deployments.apps
    - statefulsets.apps
  definitionRef:
    name: metricstrait.standard.oam.dev
  workloadRefPath: spec.workloadRef
  extension:
    template: |-
      #metrics: {
        // +usage=format of the metrics, default as prometheus
        // +short=f
        format: *"prometheus" | string
        path: */metrics" | string
```

`definition.yaml`

# Exercise 2



KubeCon



CloudNativeCon

North America 2020

*Virtual*

## Add a new capability to KubeVela

1. Create a slack bot [https://api.slack.com/apps?new\\_app=1](https://api.slack.com/apps?new_app=1)
2. `vela cap center config mycap https://github.com/oam-dev/catalog/tree/master/registry`
3. `vela cap ls`
4. `vela cap add mycap/kubewatch`
5. `vela comp deploy mycomp -t webservice --image crccheck/hello-world --port 8000 --app myapp`
6. `vela kubewatch mycomp --app myapp --webhook https://hooks.slack.com/<yourid>`



# KubeVela in Nutshell



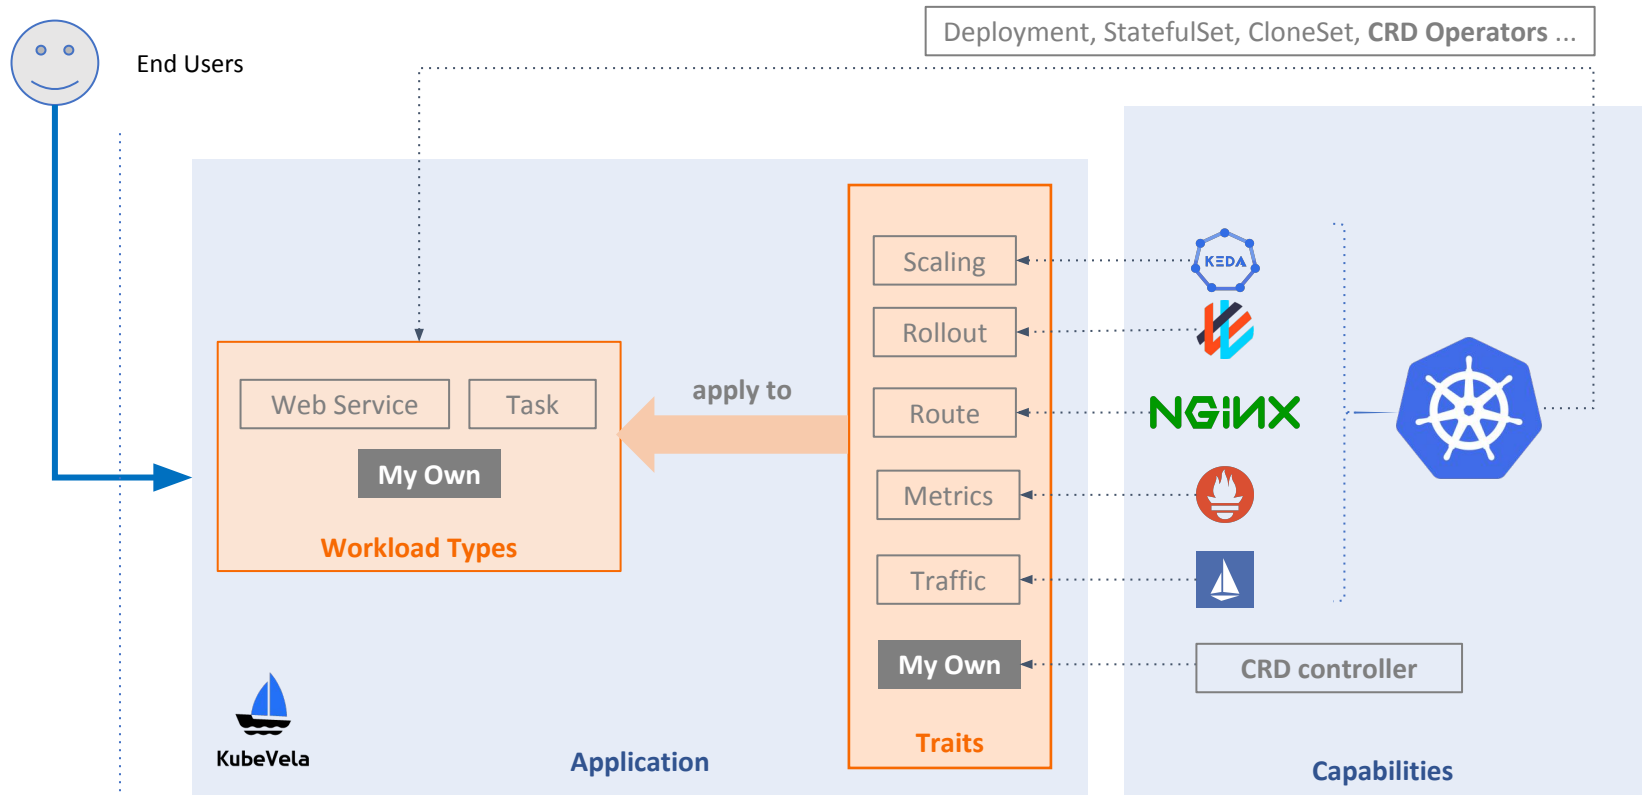
KubeCon



CloudNativeCon

North America 2020

*Virtual*



Extensible UI

"Application" as the main API

Capability discovery and management (COA)

# Overall Architecture



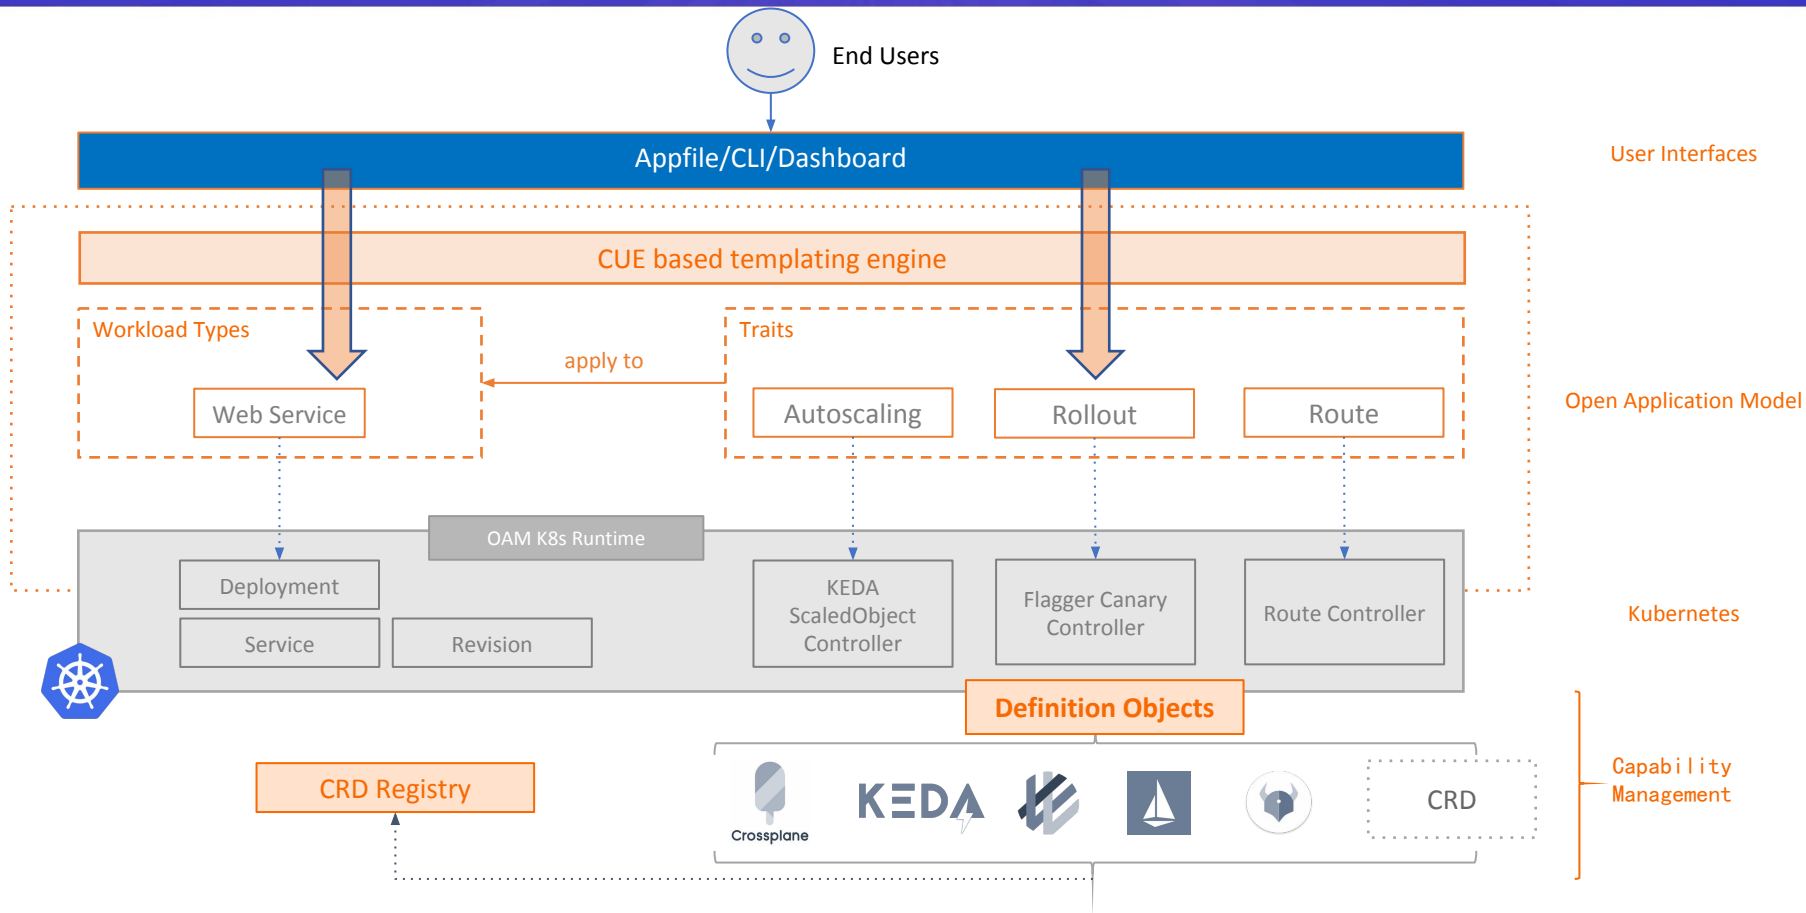
KubeCon



CloudNativeCon

North America 2020

Virtual



# What is Crossplane?



North America 2020

*Virtual*

- Handles the **infrastructure** for your applications
- Based on Kubernetes control plane (CRDs and controllers)
- CNCF Sandbox project in June 2020, from the creators of Rook (CNCF graduated)
- 3 main feature areas
  - Provision infrastructure declaratively using the Kubernetes API
  - Build your own declarative infrastructure API without code
  - Run and deploy applications alongside infrastructure



# Exercise 3



KubeCon



CloudNativeCon

North America 2020

*Virtual*

1. Prepare cloud resources with Crossplane


<https://crossplane.io/docs/v0.13/getting-started/install-configure.htm>

2. Configure the application to consume cloud resources
  - a. Create RDS workload definition
  - b. Deploy comps running with these definitions



## 1) Provision Infrastructure with K8s API

- Cloud and on-prem infrastructure is represented as a CRD
- Declaratively configured (e.g. YAML)
- Controllers reconcile infrastructure CRDs with cloud provider or on-prem APIs (e.g., GCP, AWS, Azure, Alibaba, Packet)
- Provision infrastructure with `kubectl` or any tool that talks with the Kubernetes API



# Deep Dive Crossplane



KubeCon



CloudNativeCon

North America 2020

Virtual

## Ex: Cloud Primitives for AWS - [https://doc.crds.dev/](https://doc.crds.dev/crossplane/provider-aws)

### crossplane/provider-aws

[github.com/crossplane/provider-aws/tree/master](https://github.com/crossplane/provider-aws/tree/master)

CRDs discovered: 39

Last parsed: Sun, 25 Oct 2020 16:00:33 +0000

Kind	Group	Version
<a href="#">Certificate</a>	acm.aws.crossplane.io	v1alpha1
<a href="#">CertificateAuthority</a>	acmpca.aws.crossplane.io	v1alpha1
<a href="#">CertificateAuthorityPermission</a>	acmpca.aws.crossplane.io	v1alpha1
<a href="#">Provider</a>	aws.crossplane.io	v1alpha3
<a href="#">ProviderConfig</a>	aws.crossplane.io	v1beta1
<a href="#">ProviderConfigUsage</a>	aws.crossplane.io	v1beta1
<a href="#">CacheCluster</a>	cache.aws.crossplane.io	v1alpha1
<a href="#">CacheSubnetGroup</a>	cache.aws.crossplane.io	v1alpha1
<a href="#">ReplicationGroup</a>	cache.aws.crossplane.io	v1beta1
<a href="#">DBSubnetGroup</a>	database.aws.crossplane.io	v1beta1
<a href="#">DynamoTable</a>	database.aws.crossplane.io	v1alpha1
<a href="#">RDSInstance</a>	database.aws.crossplane.io	v1beta1

Kubernetes Clusters

Databases

Networking

Certificates

Caches

Message Queues

....

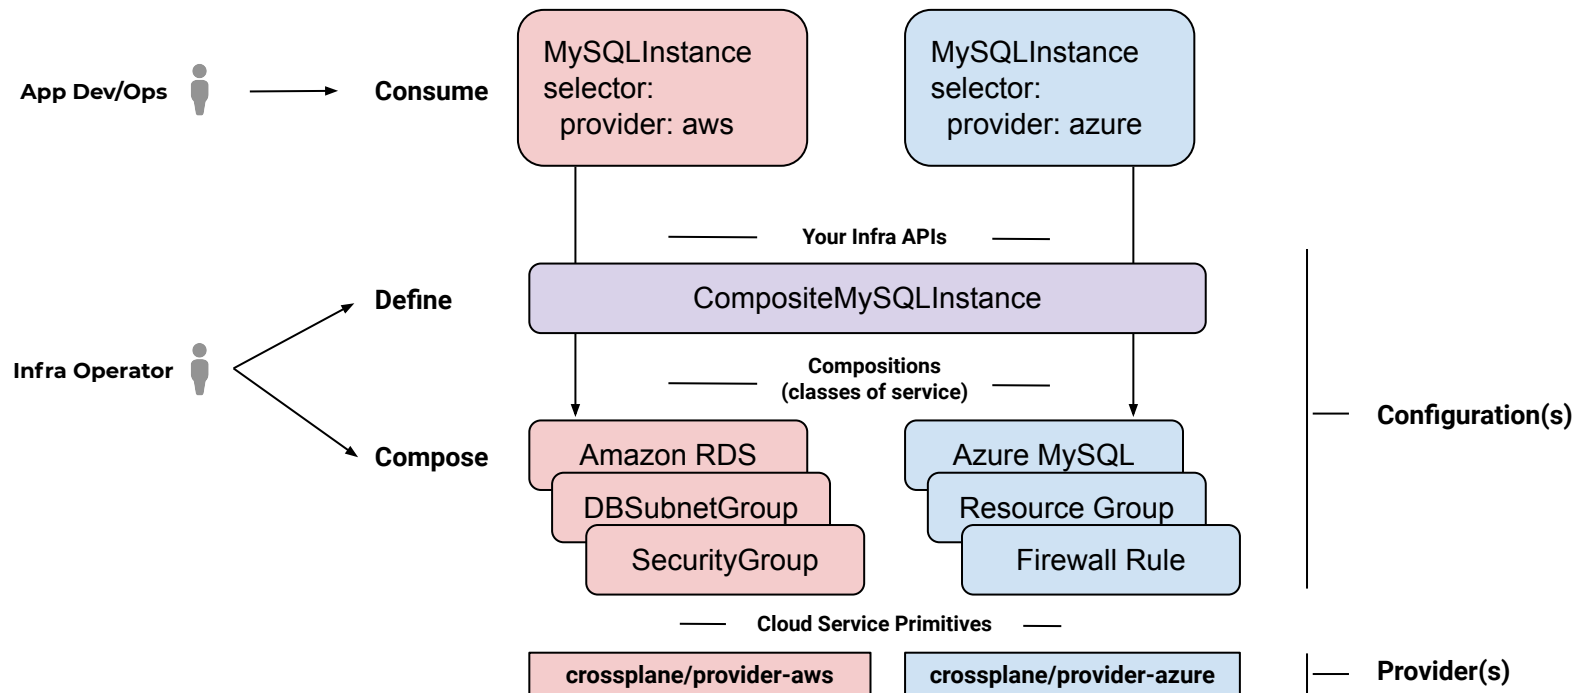


## 2) Offer declarative Infra APIs for teams to use

- Define, compose, and offer your own infrastructure API abstractions on top of the cloud service primitives included
- Example: **Define** a MySQL resource that is **composed** of Azure MySQL, resource group, and firewall rule. **Offer** it for applications to consume on demand.
- Hide infrastructure complexity and include policy guardrails
- **No code** required, it's all **declarative**



## Composition - define & compose your own Infra APIs







## 3) Run and deploy applications

- Support for Open Application Model (OAM)
- Deploy applications alongside infrastructure
  - Standardize on a single workflow for both
- Strong “separation of concerns”
  - Infrastructure operators - provide infrastructure and services
  - Application developers - build application components independent of infrastructure
  - Application operators - compose, deploy, and run application configurations



# Crossplane Arch Summary




KubeCon



CloudNativeCon

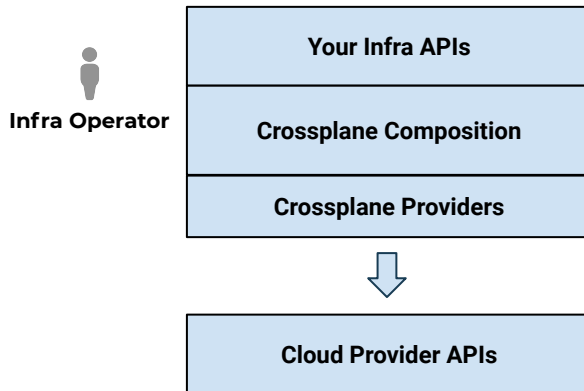
North America 2020

*Virtual*



Consume Infra APIs

Provide Infra APIs



- Open Application Model (OAM)
- k8s core resources (Deployments, Services, ...)

- Abstraction - define your own claim kinds and schema
- Offer multiple classes of service for an Infra API
- Compose cloud service primitives
- Bridge cloud APIs into k8s - cloud service primitives



# Exercise 4



KubeCon



CloudNativeCon

North America 2020

*Virtual*

## Build and offer an infrastructure API to your team

1. **Define** your platform's composite resources and claims
  - a. <https://github.com/upbound/platform-ref-aws/>
  - b. Network, Cluster, PostgreSQL
2. **Push** this configuration to a registry
3. **Install** the configuration into a Crossplane instance
4. **Offer** the desired claims to your team(s)
5. Teams **provision** their infrastructure self-service, with the policy/config you declared



# Community - OAM/KubeVela



KubeCon



CloudNativeCon

North America 2020

*Virtual*

- **OAM Specification**

- Now [v0.2.2 working draft](#), moving to beta release with full backward compatibility

- **KubeVela**


- *Developer Preview* stage with features still WIP, **NOT** ready for production.
  - Roadmap: <https://github.com/oam-dev/kubevela/projects/1>
    - v1.0.0 release targets at Dec. 2020
  - Current feature set:
    - *Appfile, CLI, dashboard (preview)*
    - **Web Service** & **Task** workload types, **Route, Rollout** (Flagger) & **Autoscaling** (KEDA)

- **Community**

- Gitter: <https://gitter.im/oam-dev/>
  - Slack: <https://cloud-native.slack.com/messages/kubevela/>

## Get involved with this CNCF Sandbox project!

- **Website:** <https://crossplane.io/>
- **Slack:** <https://slack.crossplane.io/>
- **Docs:** <https://crossplane.io/docs>
- **GitHub:** <https://github.com/crossplane/crossplane>
- **Blog:** <https://blog.crossplane.io/>
- **Twitter:** [https://twitter.com/crossplane\\_io](https://twitter.com/crossplane_io)
- **Youtube:** [Crossplane Youtube](#)



## Working towards (and beyond) a v1.0 release for end of year!

- **Crossplane Core**

- Composition - composition revisions, validation, custom compositions
- Package Manager - dependency resolution, multiple versions
- Hardening, metrics, APIs to v1beta1+

- **Providers for AWS, GCP, Azure, Alibaba**

- Rapidly expanding cloud services to 90% coverage
- AWS ACK & Azure ASO code-gen of Crossplane providers
- Wrapping stateless Terraform providers for clouds w/o code-gen

- **Open Application Model (OAM)**

- oam-kubernetes-runtime APIs to v1beta1 & subchart install
- health scopes & additional trait support



